

# HelpMate AI Project\_V3

September 2, 2025

```
[15]: !pip install sentence-transformers transformers torch diskcache pypdf ↵  
      ↪scikit-learn
```

```
import os, re, hashlib, warnings  
from typing import List, Tuple, Dict  
from pypdf import PdfReader  
from sentence_transformers import SentenceTransformer, CrossEncoder  
from transformers import pipeline  
from diskcache import Cache  
from sklearn.metrics.pairwise import cosine_similarity  
import numpy as np
```

Requirement already satisfied: sentence-transformers in  
/usr/local/lib/python3.12/dist-packages (5.1.0)  
Requirement already satisfied: transformers in /usr/local/lib/python3.12/dist-  
packages (4.55.4)  
Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-packages  
(2.8.0+cu126)  
Requirement already satisfied: diskcache in /usr/local/lib/python3.12/dist-  
packages (5.6.3)  
Requirement already satisfied: pypdf in /usr/local/lib/python3.12/dist-packages  
(6.0.0)  
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.12/dist-  
packages (1.6.1)  
Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-packages  
(from sentence-transformers) (4.67.1)  
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages  
(from sentence-transformers) (1.16.1)  
Requirement already satisfied: huggingface-hub>=0.20.0 in  
/usr/local/lib/python3.12/dist-packages (from sentence-transformers) (0.34.4)  
Requirement already satisfied: Pillow in /usr/local/lib/python3.12/dist-packages  
(from sentence-transformers) (11.3.0)  
Requirement already satisfied: typing\_extensions>=4.5.0 in  
/usr/local/lib/python3.12/dist-packages (from sentence-transformers) (4.15.0)  
Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-  
packages (from transformers) (3.19.1)  
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.12/dist-

packages (from transformers) (2.0.2)  
 Requirement already satisfied: packaging>=20.0 in  
 /usr/local/lib/python3.12/dist-packages (from transformers) (25.0)  
 Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.12/dist-  
 packages (from transformers) (6.0.2)  
 Requirement already satisfied: regex!=2019.12.17 in  
 /usr/local/lib/python3.12/dist-packages (from transformers) (2024.11.6)  
 Requirement already satisfied: requests in /usr/local/lib/python3.12/dist-  
 packages (from transformers) (2.32.4)  
 Requirement already satisfied: tokenizers<0.22,>=0.21 in  
 /usr/local/lib/python3.12/dist-packages (from transformers) (0.21.4)  
 Requirement already satisfied: safetensors>=0.4.3 in  
 /usr/local/lib/python3.12/dist-packages (from transformers) (0.6.2)  
 Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-  
 packages (from torch) (75.2.0)  
 Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-  
 packages (from torch) (1.13.3)  
 Requirement already satisfied: networkx in /usr/local/lib/python3.12/dist-  
 packages (from torch) (3.5)  
 Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages  
 (from torch) (3.1.6)  
 Requirement already satisfied: fsspec in /usr/local/lib/python3.12/dist-packages  
 (from torch) (2025.3.0)  
 Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in  
 /usr/local/lib/python3.12/dist-packages (from torch) (12.6.77)  
 Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in  
 /usr/local/lib/python3.12/dist-packages (from torch) (12.6.77)  
 Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in  
 /usr/local/lib/python3.12/dist-packages (from torch) (12.6.80)  
 Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in  
 /usr/local/lib/python3.12/dist-packages (from torch) (9.10.2.21)  
 Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in  
 /usr/local/lib/python3.12/dist-packages (from torch) (12.6.4.1)  
 Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in  
 /usr/local/lib/python3.12/dist-packages (from torch) (11.3.0.4)  
 Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in  
 /usr/local/lib/python3.12/dist-packages (from torch) (10.3.7.77)  
 Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in  
 /usr/local/lib/python3.12/dist-packages (from torch) (11.7.1.2)  
 Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in  
 /usr/local/lib/python3.12/dist-packages (from torch) (12.5.4.2)  
 Requirement already satisfied: nvidia-cusparselt-cu12==0.7.1 in  
 /usr/local/lib/python3.12/dist-packages (from torch) (0.7.1)  
 Requirement already satisfied: nvidia-nccl-cu12==2.27.3 in  
 /usr/local/lib/python3.12/dist-packages (from torch) (2.27.3)  
 Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in  
 /usr/local/lib/python3.12/dist-packages (from torch) (12.6.77)  
 Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in

```

/usr/local/lib/python3.12/dist-packages (from torch) (12.6.85)
Requirement already satisfied: nvidia-cublas-cu12==1.11.1.6 in
/usr/local/lib/python3.12/dist-packages (from torch) (1.11.1.6)
Requirement already satisfied: triton==3.4.0 in /usr/local/lib/python3.12/dist-
packages (from torch) (3.4.0)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-
packages (from scikit-learn) (1.5.1)
Requirement already satisfied: threadpoolctl>=3.1.0 in
/usr/local/lib/python3.12/dist-packages (from scikit-learn) (3.6.0)
Requirement already satisfied: hf-xet<2.0.0,>=1.1.3 in
/usr/local/lib/python3.12/dist-packages (from huggingface-hub>=0.20.0->sentence-
transformers) (1.1.8)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in
/usr/local/lib/python3.12/dist-packages (from sympy>=1.13.3->torch) (1.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in
/usr/local/lib/python3.12/dist-packages (from jinja2->torch) (3.0.2)
Requirement already satisfied: charset_normalizer<4,>=2 in
/usr/local/lib/python3.12/dist-packages (from requests->transformers) (3.4.3)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.12/dist-
packages (from requests->transformers) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.12/dist-packages (from requests->transformers) (2.5.0)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.12/dist-packages (from requests->transformers) (2025.8.3)

```

## 1 New Section

```

[16]: # ----- CONFIG -----
PDF_PATH = "/content/sample_data/Principal-Sample-Life-Insurance-Policy.pdf"
CACHE_DIR = "/content/sample_data/cache"
#EMBED_MODEL = "all-MiniLM-L6-v2"
EMBED_MODEL = "multi-qa-MiniLM-L6-cos-v1"
RERANK_MODEL = "cross-encoder/ms-marco-MiniLM-L-6-v2"
LLM_MODEL = "google/flan-t5-base" # lightweight; replace with "gpt2" or
↳ OpenAI API

TOPK_RETRIEVE = 15 # recall first
TOPK_RERANK = 3 # precision later
#CHUNK_WORDS = 180
#CHUNK_OVERLAP = 40
CHUNK_WORDS = 300
CHUNK_OVERLAP = 50
RERANK_CLIP_WORDS = 100 # Reduced from 180

warnings.filterwarnings("ignore", category=UserWarning)

```

```

[17]: # ----- HELPERS -----
def clean_text(text: str) -> str:
    return re.sub(r'\s+', ' ', (text or "")).strip()

def sentence_split(text: str) -> List[str]:
    # Simple sentence split; robust enough for policy docs
    parts = re.split(r'(?<=[\.\?!\])\s+', text)
    return [s.strip() for s in parts if s.strip()]

def chunk_page_sentences(sentences: List[str], max_words=CHUNK_WORDS,
    overlap=CHUNK_OVERLAP) -> List[str]:
    chunks = []
    curr, count = [], 0
    for sent in sentences:
        w = len(sent.split())
        if count + w > max_words and curr:
            chunks.append(" ".join(curr))
            # create overlap
            if overlap > 0:
                back = []
                words_kept = 0
                for s in reversed(curr):
                    sw = len(s.split())
                    if words_kept + sw >= overlap:
                        back.insert(0, s)
                        break
                words_kept += sw
                back.insert(0, s)
            curr = back
            count = sum(len(s.split()) for s in curr)
        else:
            curr, count = [], 0
        curr.append(sent)
        count += w
    if curr:
        chunks.append(" ".join(curr))
    return chunks

def load_pdf_chunks(pdf_path: str, max_words=CHUNK_WORDS,
    overlap=CHUNK_OVERLAP) -> Tuple[List[str], List[Dict]]:
    reader = PdfReader(pdf_path)
    chunks, metas = [], []
    for i, page in enumerate(reader.pages):
        text = clean_text(page.extract_text())
        if not text:
            continue
        sents = sentence_split(text)

```

```

        page_chunks = chunk_page_sentences(sents, max_words=max_words,
↪overlap=overlap)
        for j, ch in enumerate(page_chunks):
            chunks.append(ch)
            metas.append({"page": i + 1, "chunk_id": j})
    return chunks, metas

```

```

[18]: # ----- EMBEDDING -----
class Embedder:
    def __init__(self, model_name=EMBED_MODEL):
        self.model = SentenceTransformer(model_name)
    def embed(self, texts):
        return self.model.encode(texts, convert_to_numpy=True)

```

```

[19]: # ----- VECTOR SEARCH (SKLEARN) -----
def build_index(embeddings):
    return np.array(embeddings)

def search_index(query_vec, index, top_k):
    sims = cosine_similarity([query_vec], index)[0]
    top_ids = sims.argsort()[-top_k:][::-1]
    return top_ids.tolist()

```

```

[20]: # ----- SEARCH + RERANK -----
cache = Cache(CACHE_DIR)
reranker = CrossEncoder(RERANK_MODEL)

def _pdf_signature(pdf_path: str) -> str:
    try:
        stat = os.stat(pdf_path)
        sig = f"{os.path.basename(pdf_path)}::{stat.st_mtime_ns}::{stat.
↪st_size}"
    except FileNotFoundError:
        sig = pdf_path
    return hashlib.md5(sig.encode()).hexdigest()

def clip_words(text: str, limit_words=RERANK_CLIP_WORDS) -> str:
    words = text.split()
    if len(words) <= limit_words:
        return text
    return " ".join(words[:limit_words])

def retrieve(query: str, embedder: Embedder, index: np.ndarray, chunks:
↪List[str], metas: List[Dict],
            top_k: int = TOPK_RETRIEVE, pdf_sig: str = "") -> List[int]:
    key = ("retr", query, top_k, pdf_sig, EMBED_MODEL)
    if key in cache:

```

```

        return cache[key]
    qvec = embedder.embed([query])[0]
    ids = search_index(qvec, index, top_k)
    cache[key] = ids
    return ids

def rerank_ids(query: str, ids: List[int], chunks: List[str], top_k: int = TOPK_RERANK) -> List[int]:
    # Shorten text to reduce truncation bias in cross-encoder
    pairs = [(query, clip_words(chunks[i])) for i in ids]
    scores = reranker.predict(pairs)
    ranked = sorted(zip(ids, scores), key=lambda x: x[1], reverse=True)
    return [i for i, _ in ranked[:top_k]]

```

```

[21]: # ----- GENERATION -----
generator = pipeline("text2text-generation", model=LLM_MODEL, device=-1 )

def build_prompt(query: str, items: List[Tuple[str, Dict]]) -> str:
    # Compose a compact, citation-friendly context
    lines = []
    for t, m in items:
        p = m.get("page", "?")
        lines.append(f"[p.{p}] {t}")
    context = "\n".join(lines)
    prompt = (
        "You are an assistant that answers ONLY using the context provided.\n"
        "Rules:\n"
        "1. Summarize the answer in 2-3 clear sentences.\n"
        "2. Do not repeat sentences or phrases.\n"
        "3. Always include page numbers as citations like [p.X].\n"
        "4. If the answer is not present, reply exactly: Not found in policy.\n"
        "\n"
        f"Context:\n{context}\n"
        f"Question: {query}\n"
        "Answer:"
    )
    return prompt

def generate_answer(query: str, chosen: List[int], chunks: List[str], metas: List[Dict]) -> str:
    items = [(chunks[i], metas[i]) for i in chosen]
    prompt = build_prompt(query, items)
    out = generator(prompt, max_new_tokens=10, do_sample=False)[0]["generated_text"]
    return out.strip()

```

Device set to use cpu

```
[22]: # ----- MAIN PIPELINE -----
def run_pipeline(pdf_path: str, queries: List[str]) -> None:
    pdf_sig = _pdf_signature(pdf_path)

    # Step 1: Ingest
    chunks, metas = load_pdf_chunks(pdf_path, CHUNK_WORDS, CHUNK_OVERLAP)
    print(f"Loaded {len(chunks)} chunks from PDF.")

    # Step 2: Embed + index
    embedder = Embedder()
    embeddings = embedder.embed(chunks)
    index = build_index(embeddings)

    # Step 3: Query loop
    for q in queries:
        # retrieve a wide set, then rerank down to the best
        cand_ids = retrieve(q, embedder, index, chunks, metas,
        ↪top_k=TOPK_RETRIEVE, pdf_sig=pdf_sig)
        best_ids = rerank_ids(q, cand_ids, chunks, top_k=TOPK_RERANK)

        # Pretty print top-3 with pages
        print("\n=====")
        print("Query:", q)
        print("\nTop Retrieved (after rerank):")
        for rank, i in enumerate(best_ids, 1):
            pg = metas[i]["page"]
            print(f"{rank}) [p.{pg}] {chunks[i][:220]}...")

        # Generate final answer
        answer = generate_answer(q, best_ids, chunks, metas)
        print("\nFinal Answer:")
        print(answer)
        print("=====")

[23]: # ----- RUN -----
if __name__ == "__main__":
    # Example queries - replace with your own
    queries = [
        "What is the grace period for premium payment?",
        "Who is eligible under this policy?",
        "List the major exclusions."
    ]
    if not os.path.exists(PDF_PATH) or os.path.isdir(PDF_PATH):
        print(" Please set PDF_PATH at the top of the script to your policy_
        ↪PDF file.")
    else:
        run_pipeline(PDF_PATH, queries)
```

Loaded 95 chunks from PDF.

modules.json: 0%| | 0.00/349 [00:00<?, ?B/s]  
config\_sentence\_transformers.json: 0%| | 0.00/116 [00:00<?, ?B/s]  
README.md: 0.00B [00:00, ?B/s]  
sentence\_bert\_config.json: 0%| | 0.00/53.0 [00:00<?, ?B/s]  
config.json: 0%| | 0.00/612 [00:00<?, ?B/s]  
model.safetensors: 0%| | 0.00/90.9M [00:00<?, ?B/s]  
tokenizer\_config.json: 0%| | 0.00/383 [00:00<?, ?B/s]  
vocab.txt: 0.00B [00:00, ?B/s]  
tokenizer.json: 0.00B [00:00, ?B/s]  
special\_tokens\_map.json: 0%| | 0.00/112 [00:00<?, ?B/s]  
config.json: 0%| | 0.00/190 [00:00<?, ?B/s]

Token indices sequence length is longer than the specified maximum sequence length for this model (907 > 512). Running this sequence through the model will result in indexing errors

=====

Query: What is the grace period for premium payment?

Top Retrieved (after rerank):

- 1) [p.20] This policy has been updated effective January 1, 2014 PART II - POLICY ADMINISTRATION GC 6004 Section B - Premiums, Page 1 Section B - Premiums Article 1 - Payment Responsibility; Due Dates; Grace Period The Policyholde...
- 2) [p.23] This policy has been updated effective January 1, 2014 PART II - POLICY ADMINISTRATION GC 6005 Section C - Policy Termination, Page 1 Section C - Policy Termination Article 1 - Failure to Pay Premium This Group Policy wi...
- 3) [p.44] A Dependent must apply for individual purchase and the first premium for the individual policy must be paid to The Principal within 31 days after the date Dependent Life Insurance for the Dependent terminates under this ...

Final Answer:

31 days.

=====

=====

Query: Who is eligible under this policy?

Top Retrieved (after rerank):

- 1) [p.26] This policy has been updated effective January 1, 2014 PART III - INDIVIDUAL REQUIREMENTS AND RIGHTS GC 6006 Section A - Eligibility, Page 1 PART III - INDIVIDUAL REQUIREMENTS AND RIGHTS Section A - Eligibility Article 1...



- 2) [p.27] This policy has been updated effective January 1, 2014 PART III - INDIVIDUAL REQUIREMENTS AND RIGHTS GC 6006 Section A - Eligibility, Page 2 If a Member's Dependent is employed and is covered under group term life covera...
- 3) [p.41] This policy has been updated effective January 1, 2014 PART III - INDIVIDUAL REQUIREMENTS AND RIGHTS GC 6010 Section E - Reinstatement, Page 2 If coverage for a Member or Dependent terminates because the person is outsid...

Final Answer:

A person will be eligible for Member Life Insurance on the date the person completes 30 consecutive days of continuous Active Work with the Policyholder as a Member. In no circumstance will a person be eligible for Member Life Insurance under this Group Policy if the person is eligible under any other Group Term Life Insurance policy underwritten by The Principal. Article 1 - Member Life Insurance A person will be eligible for Member Life Insurance on the date the person completes 30 consecutive days of continuous Active Work with the Policyholder as a Member. In no circumstance will a person be eligible for Member Life Insurance under this Group Policy if the person is eligible under any other Group Term Life Insurance policy underwritten by The Principal. Article 2 - Member Accidental Death and Dismemberment Insurance A person will be eligible for Member Accidental Death and Dismemberment Insurance on the latest of:

=====

=====

Query: List the major exclusions.

Top Retrieved (after rerank):

- 1) [p.14] Not included are rest homes, homes for the aged, nursing homes, or places for treatment of mental disease, drug addiction, or alcoholism. Terminally Ill A Member will be considered Terminally Ill, for Accelerated Benefit...
- 2) [p.23] fails to maintain the participation percentages requirements of PART II, Section A with respect to eligible employees, excluding those for whom Proof of Good Health is not satisfactory to The Principal; or c. fails to ma...
- 3) [p.10] This policy has been updated effective January 1, 2014 GC 6002 PART I - DEFINITIONS, PAGE 2 The legally recognized union of two eligible individuals of the same sex established according to law. Civil Union Partner For t...

Final Answer:

Rest homes, homes for the aged, nursing homes, or places for treatment of mental disease, drug addiction, or alcoholism.

=====

[23]: